

← Back to Sample Reports

SANITIZED REAL-REPOSITORY EXAMPLE

Evaluated source repository: Kestrel / AI Service Call Agent · commit 5e65843

Kestrel — Agentic AI Assurance Report

A real-repository assurance example showing what the AI implementation can reach and cause, where authority evidence ends, and why the bounded Assurance Decision is REVIEW.

ASSURANCE DECISION

REVIEW

Bounded decision within evaluated scope and available evidence.

DECISION BASIS

HAIEC established **44 consequential action paths** the implementation is capable of reaching. Code capability is established for these paths, while policy authorization, effective runtime permission and observed execution were not established by the available evidence.

What this means: the implementation has meaningful action capability. The remaining assurance question is whether those capabilities are authorized, effectively permitted at runtime, and actually exercised as intended.

What to do next: prioritize the highest-consequence paths and establish policy, runtime permission and observation evidence before making a stronger assurance claim.

 Explore System Constellation

 Download / Evidence ▾



HAIEC
HUMAN-AI INTERDEPENDENT ECOSYSTEM

Selected from the 44 evaluated action paths. Five planes are independent evidence states — NOT ASSESSED is a failure.



BUSINESS ACTION

TENANT TIMEZONE CONFIG READ

handle_create_booking → calendar_service.py
· ai_agent_configs table query – reads tenant
timezone configuration

Requested	○ NOT ASSESSED
Policy Authorized	○ NOT ASSESSED
Effectively Granted	○ NOT ASSESSED
Code Capable	● ESTABLISHED
Observed	○ NOT ASSESSED

[Show in System Constellation →](#)

BUSINESS ACTION

TENANT LOCATIONS READ

handle_create_booking → template_schema.py
· locations table query – reads active tenant
locations

Requested	○ NOT ASSESSED
Policy Authorized	○ NOT ASSESSED
Effectively Granted	○ NOT ASSESSED
Code Capable	● ESTABLISHED
Observed	○ NOT ASSESSED

[Show in System Constellation →](#)

BUSINESS ACTION

AI CONFIG TEMPLATE READ

handle_create_booking → template_schema.py
· ai_config_templates table query – reads age
nt configuration template data

Requested	○ NOT ASSESSED
Policy Authorized	○ NOT ASSESSED
Effectively Granted	○ NOT ASSESSED
Code Capable	● ESTABLISHED
Observed	○ NOT ASSESSED

[Show in System Constellation →](#)

BUSINESS ACTION

CUSTOMER MESSAGE RECORDS READ

handle_create_booking → transactional_sms.p
y · sms_messages table query – reads recent m
essage records for rate limiting

Requested	○ NOT ASSESSED
Policy Authorized	○ NOT ASSESSED
Effectively Granted	○ NOT ASSESSED
Code Capable	● ESTABLISHED
Observed	○ NOT ASSESSED

[Show in System Constellation →](#)

[Show all 44 evaluated action paths](#)

⚠ Evidence coverage is partial. Bounded and unknown states below are accurate — they are not failures and not passes.

Full Assurance Report

Executive

Technical

Auditor



Assurance Report

HAIEC traces consequential AI action paths and compares code capability with the authority and evidence available for the same evaluated scope.

AI System: Kestrel — AI Service Call Agent

AI System ID:

Evaluation ID:

Evaluation snapshot: 2026-09-14T08:38:17.814Z

Repository commit: not bound

Methodology: 1.1

Evidence Bundle Digest: 3e21e47c0fdfa89d...

ASSURANCE DECISION

REVIEW

REVIEW — bounded to the evaluated scope and available evidence.

No qualifying evidence reference is bound to this statement.

EXECUTIVE SUMMARY

HAIEC evaluated this system and reached REVIEW. 44 path(s) not assessed on requested; 44 path(s) not assessed on policyAuthorized; 44 path(s) not assessed on effectivelyGranted; 44 path(s) established on codeCapable; 44 path(s) not assessed on observed

- 44 consequential action path(s) were evaluated; 44 have established code capability (what the code can reach or cause).
- 44 path(s) have established code capability without established authority evidence — policy authorization, effective permission, and observed execution remain separately stated, never inferred.
- 1322 supporting evidence trace(s) document how these paths were reconstructed.
- 31 canonical context-binding relation(s) were established; see Tenant & Subject Binding for whose context controls each action.
- 555 evidence frontier item(s) remain — see Evidence Coverage & Frontier.
- Next: Capture runtime evidence for action path 019bb0db3c13f0dd... before claiming observation.

No qualifying evidence reference is bound to this statement.

DECISION BASIS

Why the canonical assurance engine reached this disposition, from the evaluated claims bound to this exact evaluation.

- Privileged action authorization — Insufficient evidence. Coverage is UNKNOWN — absence of findings cannot become proof of absence.
- Secret exposure prevention — Insufficient evidence. Coverage is UNKNOWN — absence of findings cannot become proof of absence.
- Tenant boundary enforcement — Insufficient evidence. Coverage is UNKNOWN — absence of findings cannot become proof of absence.
- Untrusted input boundary — Insufficient evidence. Coverage is UNKNOWN — absence of findings cannot become proof of absence.
- Ai inventory completeness — Not assessed. No qualifying evaluation of this claim occurred.
- Evidence coverage was unknown for an evaluated claim.
- A claim was not evaluated.

► Inspect proof (1 reference)

ACTION ASSURANCE

Availability: partial. 6 frontier surface(s).

Limitations: CREDENTIAL_EFFECTIVE_GRANT: Credential attached but no source-qualified use — EFFECTIVELY_GRANTED = NOT_ASSESSED

No qualifying evidence reference is bound to this statement.

TENANT & SUBJECT BINDING

31 canonical context-binding relation(s) across 23 evaluated path(s). Tenant context detected on 9 relation(s); tenant restriction bound to action on 30; tenant tied to authenticated subject on 0; 31 relation(s) remain partial or unknown.

- Path 084ee19768725f5a... (TENANT_CONTEXT): tenant context detected; tenant restriction bound to action; tenant not tied to authenticated subject, value origin unknown — relation ad3_rel_kga9l8..., state partial.
- Path 084ee19768725f5a... (TENANT_CONTEXT): no tenant context detected; tenant restriction bound to action; tenant not tied to authenticated subject, value origin unknown, authentication auth not found — relation ad3_rel_utdz02..., state partial.
- Path 2cf2e1fd31e82f6d... (TENANT_CONTEXT): no tenant context detected; tenant restriction bound to action; tenant not tied to authenticated subject, value origin unknown, authentication auth not found — relation ad3_rel_7f6mte..., state partial.
- Path 2f6204531b481b20... (TENANT_CONTEXT): no tenant context detected; tenant restriction bound to action; tenant not tied to authenticated subject, value origin unknown, authentication auth not found — relation ad3_rel_qnq2wx..., state partial.

► Inspect proof (3 references)

MATERIAL MISMATCH / CONSEQUENCE

44 evaluated path(s) have established code capability without established authority evidence.
5 representative path(s) shown.

- Path 019bb0db3c13f0dd...: code capability is established, while effective permission has not been established from the evaluated evidence and policy authorization has not been established from the evaluated evidence.
- Path 084ee19768725f5a...: code capability is established, while effective permission has not been established from the evaluated evidence and policy authorization has not been established from the evaluated evidence.
- Path 1e33aa235efc0f17...: code capability is established, while effective permission has not been established from the evaluated evidence and policy authorization has not been established from the evaluated evidence.
- Path 2cf2e1fd31e82f6d...: code capability is established, while effective permission has not been established from the evaluated evidence and policy authorization has not been established from the evaluated evidence.
- Path 2f6204531b481b20...: code capability is established, while effective permission has not been established from the evaluated evidence and policy authorization has not been established from the evaluated evidence.
- + 39 additional path(s) not shown in this profile.

► Inspect proof (5 references)

REQUIRED EVIDENCE ACTIONS

Evidence-closure recommendations derived deterministically from missing or unsupported evidence states. These are report-derived recommendations, not canonical policy obligations.

- Capture runtime evidence for action path 019bb0db3c13f0dd... before claiming observation.
- Capture runtime evidence for action path 084ee19768725f5a... before claiming observation.
- Capture runtime evidence for action path 1e33aa235efc0f17... before claiming observation.
- Capture runtime evidence for action path 2cf2e1fd31e82f6d... before claiming observation.
- Capture runtime evidence for action path 2f6204531b481b20... before claiming observation.
- + 128 additional evidence-closure action(s) not shown in this profile.

No qualifying evidence reference is bound to this statement.

EVIDENCE COVERAGE & FRONTIER

555 frontier item(s): 554 partial, 1 unknown. 5 decision-relevant item(s) shown for this profile; 550 additional frontier item(s) remain in the full Technical/Auditor report.

- CONSEQUENCE_COMMIT_MODEL:handler-op:tool_impl_659:SINK:pydb_hvac_agent/app/services/calendar_service.py:50:22 — Partially established (REASON_NOT_ESTABLISHED)
- CONSEQUENCE_COMMIT_MODEL:handler-op:tool_impl_659:SINK:pydb_hvac_agent/app/services/template_schema.py:1140:18 — Partially established (REASON_NOT_ESTABLISHED)
- CONSEQUENCE_COMMIT_MODEL:handler-op:tool_impl_659:SINK:pydb_hvac_agent/app/services/template_schema.py:1294:18 — Partially established (REASON_NOT_ESTABLISHED)
- CONSEQUENCE_COMMIT_MODEL:handler-op:tool_impl_659:SINK:pydb_hvac_agent/app/services/transactional_sms.py:190:21 — Partially established (REASON_NOT_ESTABLISHED)
- CONSEQUENCE_COMMIT_MODEL:handler-op:tool_impl_659:SINK:pydb_hvac_agent/app/utils/lifecycle_tracker.py:117:9 — Partially established (REASON_NOT_ESTABLISHED)

Limitations: Evidence frontier is a deterministic projection of unestablished coverage; it is not a risk score.; Unknown != failure; SOURCE_GAP != risk.

► Inspect proof (1336 references)

CONSEQUENTIAL ACTION PATHS

5 of 44 evaluated consequential action path(s) shown. Selected for presentation; not an exhaustive absence claim.

► Inspect proof (5 references)

SYSTEM CONSTELLATION

A same-evaluation System Constellation projection is bound to this report for causal context. Exact path-level navigation is not available in this report version.

No qualifying evidence reference is bound to this statement.

CAPABILITY & AUTHORITY ALIGNMENT

Showing 5 of 44 evaluated path(s). Selected for presentation; not an exhaustive absence claim.

Path 2f6204531b481b20...

requested

Not assessed

policy Authorized

Not assessed

effectively Granted

Not assessed

code Capable

Established

observed

Not assessed

Path 019bb0db3c13f0dd...

requested

Not assessed

policy Authorized

Not assessed

effectively Granted

Not assessed

code Capable

Established

observed

Not assessed

Path 084ee19768725f5a...

requested

Not assessed

policy Authorized

Not assessed

effectively Granted

Not assessed

code Capable

Established

observed

Not assessed

Path 1e33aa235efc0f17...

requested

Not assessed

policy Authorized

Not assessed

effectively Granted

Not assessed

code Capable

Established



Path 2cf2e1fd31e82f6d...

requested

Not assessed

policy Authorized

Not assessed

effectively Granted

Not assessed

code Capable

Established

observed

Not assessed

This report is projected from the exact Assurance Evidence Bundle. It does not recompute the assurance decision, mutate canonical evidence states, or rerun analyzers.

Generated by HAIEC Evidence-Bound Assurance · Evaluation ref HAIEC-KESTREL-EVAL-5e65843 · Snapshot 2026-09-14T08:38:17.814Z · Source commit 5e65843fddfe5f907485b798e464ad37b3b3b2c7



High assurance in every consequence.

Evidence-Bound Assurance for consequential AI systems.

[Explore Sample AI System →](#)

[Scope an Enterprise POC](#)

PRODUCT

Platform

System Constellation

SOLUTIONS

Enterprise POC

Assurance Firms



HAIEC
HUMAN AI INTERFERENCE EVALUATION COMPANY

Runtime Security Testing

[AI Inventory](#)

[Evidence](#)

[Assurance](#)

— **DEVELOPERS**

[Developers](#)

[Developer Security](#)

[AI AppSec](#)

[MCP Tenant Isolation](#)

[LLMVerify](#)

[Packages](#)

— **RESOURCES**

[Learn](#)

[Resources](#)

[What Is HAIEC](#)

[How It Works](#)

[Methodology](#)

[Research](#)

[Regulations](#)

[Frameworks](#)

[Sample AI System](#)

[Sample Reports](#)

[Blog](#)

[Documentation](#)

[NYC Local Law 144](#)

[Colorado AI Act](#)

Home

About

Contact

Trust

Privacy

Terms

Pricing

HAIEC Updates

Product updates, research, technical assurance insights and material regulatory developments.

Subscribe

By subscribing, you agree to our [Privacy Policy](#). Unsubscribe anytime.



Evidence-bound assurance for consequential AI systems. See what AI systems and agents can reach, what controls apply, and what the evidence actually establishes.



HAIEC
HUMAN AI EVIDENCE COMPANY

